

Basics for C++

Picked from Lippman 5th edition

Arrays

- Arrays have a fixed size
- Array declarator has the form `a[d]` for type `a` and dimension `d`. The dimension must be known at compile time, hence a `constexpr`.
- Values inside the array are default initialized:
 1. If outside a function, set to '0'
 2. Inside a function, uninitialized
- Cannot use `auto` to deduce type of the array
- No arrays of references
- Can do list initialization, which also allows skipping the dimension `d`.
- Character arrays can be initialized thorough string literals. Need to take into account the space for a null.
- Cannot copy initialize or assign arrays.
- To define a pointer or reference to an array, you need to use brackets: `int (*PArray)[10] = &arr`.
- Read decarations inside to out.
- Can use a ranged for loop on arrays.
- Compiler does not do bound checking on `[]` operand value.
- Arrays usually devolve to pointers.
- The name of the array `arr` becomes synonymous to `&arr[0]`.
- This means when you use `auto p = arr`; where `int[10] arr`, the deduced type is `int *`.
- `decltype(arr)` returns `int[10]` faithfully.
- Pointers can act as iterators on arrays.
- You can use `begin` and `end` on arrays to get 'iterators' just like vectors.
- Pointers hence support dereference, increment, comparisons etc-.
- `int *ap = ap2+4` does not add 4 to the value, but adds `4*sizeof(*ap)`, essentially it is like `&ap[4]`.
- Difference of pointers follows the same logic.
- Adding 0 to null pointers or subtracting two nul pointers is okay.
- `*ia + 4` does the dereference first.
- The builtin subscript operator on arrays allows for `ia[-2]` kind of expressions.
- C style strings are a convention of using character arrays that are null terminated.
- Doing comparisons on C-style strings ends up comparing the pointers.
- Try using `strcmp(s1, s2) < 0` or something.
- While using `strcpy`, make sure there is enough space first.
- We can use `.c_str` on string to get a C-style string, but it has type `const`

`char *.`

- Arrays of arrays are multi-dimensional arrays in C++.
- `int ia[3][4]` is array of size 3, each element is array of size 4, each element is `int`.
- We can define references to array, so we can write a statement like `int (&row)[4] = ia[1]` to get an alias for the second row.
- To use a ranged for loop, use code like:

```
for (auto &row : ia) {  
    for (auto col : row) {  
        col = cnt;  
        ++cnt;  
    }  
}
```

- We used a reference since otherwise the `auto` would make a pointer out of it.
- Try using `begin` and `end` instead.
- An array typedef looks like `typedef int intarray[4]`.

Ranged for loop

- Anything can go into after the colon `:`, that defines `begin` and `end`.
- This allows for braced initializer list, arrays, and vectors etc-.

Type conversions

- Implicit conversions happen without programmer knowledge, and occur at:
 1. Different integral types are promoted to a common large integral type in an expression.
 2. In conditions, non-bool are converted to bool
 3. In initializations, the initializer is converted to the type of the lhs.
 4. For operands of mixed types on arithmetic operators, they are converted to a common type.
 5. During function calls.
- Arithmetic conversions usually happen in two steps, first to integral level and then higher. Unsigned vs signed is often machine dependent.
- Array to pointer conversions are frequent wherever demanded. Not done only when `decltype` is used.
- 0 and `nullptr` can convert to any pointer type, a pointer to nonconst can convert to `void*` and any pointer can convert to `const void *`. Inheritance also allows for a few.
- Pointer to nonconst can convert into pointer to const.
- Classes themselves can selectively define some conversions.
- Type casts give a way to do explicit type conversions, could be `static_cast`, `const_cast` or `reinterpret_cast`

Functions

- A function is a block of code with a name, usually called using the call operator `()`.
- During a function call, function parameters are initialized and then control is transferred.
- Execution ends when a return type statement is found
- Arguments provide values for the initialization of the parameters.
- An empty parameter list still needs to be specified as `()`.
- Might instead say `void f(void)`.
- Need to specify type of each parameter separately.
- Unused parameters can be left nameless.
- The return type of a function cannot be an array or a function, but it could be a pointer to those.
- Scope: part of program text where name is visible, lifetime: part of execution that object exists.
- Parameters and other local variables are automatic objects, meaning they lose identity after block scope is over.
- To have an object whose lifetime continues across calls to the function, define it as `static`. Initialized before the first time it is defined. Eg:

```
size_t count_calls() {  
    static size_t ctr = 0;  
    return ++ctr;  
}
```

- A function declaration is definition minus the body. Also called function prototype.
- Declare the functions in header and define them in a separate source code file that imports this header.
- To enable separate compilation, provide `-c` flag which gives `.o` as output. Link the object to files together to get `./a.out`.
- Arguments can be passed by reference by declaring the parameter as a reference type.
- Pointers give indirect access to shared data.
- It is inefficient and sometimes impossible to copy objects, hence it is safer to use references.
- Unchanged reference parameters should be declared `const`.
- References can also be used to return additional information to the calling environment.
- Just as in normal initialization, top-level `const`s are ignored between parameter and argument.
- This means that the following functions cannot be distinguished so the second declaration is in error:

```
void fcn(const int i);  
void fcn(int i);
```

- The basic general rules for initialization are:
 1. ignore top level consts
 2. can use nonconst argument for low level const parameter, but not vice versa
 3. plain reference must be initialized by object of exact same type (cannot pass literals, expressions, or even const references)
- So always try to make references const if possible.
- Arrays interfere in two ways: you cannot copy them and they usually devolve into pointers.
- So the following three are the same:

```
void print(const int *);
void print(const int []);
void print(const int [10]);
```

- This means size of the array is unknown to function, specified through one of three ways:
 1. Special marker after last object of array.
 2. Sentinel pointer also passed.
 3. Explicit size parameter passed.
- The same discussion of references applies to pointers, since pointers to const can accept nonconst, but not vice-versa.
- To define an array reference parameter, we can use `void print(int (&arr)[10])`, where the bracket is important.
- To pass a multidimensional array, we would have to essentially pass a pointer to the first element of the outermost array, which means you need to know the dimensions of all the rest of the arrays.
- Command line parameters are passed to main through the “explicit size parameter” approach `int main(int argc, char **argv)`.
- To pass varying number of arguments of the same type `a`, you can use `initializer_list<a>`. It provides functionalities to assign the list, get the begin and end iterators (for range based for loops) and getting the size.
- When you want to pass multiple values to a `initializer_list` type parameter, enclose them in curly braces.
- An example:

```
void err_msg(initializer_list<string> il) {
    for (auto &s: il) {
        cout << s << " ";
    }
    cout << endl;
}
string s = "bye";
err_msg({"hello", s});
```

- We can instead use ellipses to interface with C code, but isn't recommended, since `varargs` objects are usually not copied correctly under C++.

- To terminate execution of a function with no return value, one can use `return;`
- Otherwise every return must return a value of the specified type, or which can be implicitly converted.
- A compiler may not always be able to check if all possible flows always end in a return. If this happens, runtime behavior is not defined.
- When a function returns a reference, it is specified by the return type in the prototype of the function.
- Never return a pointer or a reference to a local object.
- We can chain together method calls left-associatively: `a.foo().bar()`.
- Calls to functions that return references are lvalue, others are rvalue. The lvalue can be placed on the lhs of an assignment: `get(vec, 0) = 'A'`.
- We can't do it if the return type is a const reference.
- One can also return a `initializer_list`, which will be used to initialize the temporary return object.
- The `main` function is allowed to terminate without a return, in which case it is assumed to return 0.
- `main` is not allowed to call itself in general.
- It is usually convenient to use typedefs when returning pointer to arrays, but the general syntax looks like `type (*function(params)) [dimension]`.
- One can also use the trailing return type for declaring function prototypes: `auto func(params) -> ret_type`. For instance, `auto foo(int i) -> int (*)[10]`.
- One can also use `decltype` to simplify some declarations.
- Functions with the same name but different parameters are said to be overloaded, but note that the `main` function cannot be overloaded.
- The functions may not differ only in their return types.
- Parameters differing only in their name or in type aliasing or top level consts are not different.
- When low-level consts come in the picture, it is clear what happens on const arguments. On nonconst arguments, the compiler prefers to call the nonconst version of the function.
- The compiler does function matching to consult which function to call.
- A basic rule is that name lookup happens before type checking. This means that names do not overload across scopes. The compiler will stick with the earliest scope where it finds a match.
- This means that even if an enclosing scope has a declaration which might match, the compiler will ignore it.
- Default arguments can be specified as initializers in the prototype of the function.
- A parameter with default arguments needs all parameters to its right to also have default arguments.
- One can specify the default argument for any parameter at only place (declaration).
- Consider adding default arguments as part of the declarations in the header.
- Local arguments may not be used as default arguments, apart from this

restriction any expression works. The argument is calculated at runtime, and even if the variable is redefined in the inner scope (where the function is called with default arguments), the outer variable is used.

- A function may be declared inline, which is a request to the compiler to expand it in-place when called.
- A function is declared as inline by putting `inline` in front of the return type.
- Similarly a function may be declared `constexpr`.
- A `constexpr` function is one that can be used in a constant expression.
- A `constexpr` function must have a single return statement, and the return type and type of each parameter must be a literal type (as in simple enough for a `constexpr` to have this type).
- The return value need not be a constant expression, and the body may contain other statements as long as they have no run time consequences.
- These functions can now be used in expressions where `constexpr` is needed, like array definitions.
- Inline and `constexpr` functions may be defined multiple times, but all must match exactly.
- Usually this means we put them in headers.
- The `assert` preprocessor macros can be used to ensure that a condition we expect to be true is actually true at runtime during development.
- Looks like an inline function call: `assert (expr)`.
- A preprocessor variable called `NDEBUG` can toggle the execution of `assert`.
- By default `NDEBUG` is not defined, so `assert` runs. By adding a flag like `-D NDEBUG` during compilation, we can turn off all `asserts`.
- The compiler also provides several local variables like `__func__` which can be used during debugging.
- To do function matching, the compiler runs a number of phases:
 1. Candidate functions: set of all functions visible (from the latest scope where any function of that name is visible).
 2. Viable functions: set of all functions which have correct number of arguments and the types can be matched, maybe through conversions.
 3. Get a best match, if one does not exist, call it ambiguous and terminate.
- A function to call matching is best if the match for each argument is no worse than for any other viable function, and there is at least one argument that has a strictly better match over each of the other viable functions.
- Argument matchings are ranked as follows:
 1. An exact match: identical, array or function to pointer degradation, or top-level consts get kicked.
 2. Match through const conversion.
 3. Match through promotion.
 4. Arithmetic or pointer conversion.
 5. Class type conversion.
- Here, const conversions mean low-level constness, promotion means the process of making everything `int` and pointer conversion means casts to stuff like `void*` or from 0 to a pointer type.

- A pointer to function is generally declared as `return_type (*func)(params)`. Notice the difference from function returning pointer to array.
- Both `a = func` and `a = &func` do the same thing.
- Both `a()` and `(*a)()` are also equivalent.
- There are no conversions between pointers to functions.
- When assigning function pointers among overloaded functions, only exact matches are considered.
- We can define functions that take pointers to functions, and return pointers to functions.
- One can use `auto+decltype` or `typedefs` to make these declarations simpler.